

Insidor

How the whole thing works, every feature, and the plan to build it — in plain English

This document explains Insidor from the ground up. No coding knowledge needed. It covers what we're building, how each piece works, what happens behind the scenes, and the order we build it in.

1 · What Insidor is

The problem

Memecoins are made out of things going viral on the internet. A funny video, a strange news story, an animal with a name — someone turns it into a coin, and if they're early they make money.

The problem is **timing**. Every tool that traders currently use — Axiom, Photon, GMGN — only shows you a coin *after* it already exists and money is already flowing into it. By the time you see it there, the people who made money have already made it. You're late, and you're often the person they sell to.

The idea

Insidor works one step earlier. Instead of watching the money, we watch **the internet**. We scan social media constantly, spot posts that are exploding right now, and show them to traders *before anybody has made a coin out of them*.

That gives the user two ways to win:

- **No coin exists yet** → they can create it themselves, in one tap, and be the first person in.
- **A coin already exists** → they can buy it in the same app, still early, because we caught the trend when it started.

Why this is defensible

Anyone can build a trading screen. What's hard to copy is **two things that grow over time**: the community of traders sharing their thinking on our narratives (they won't leave, because their history and reputation is here), and our own private record of "which viral posts became coins that made money" — a dataset that gets more valuable every single day it runs, and that a new competitor starts with zero of.

2 · How it works, end to end

Here is the full journey, from someone posting a video to a trader making a trade.

1. **COLLECT** Every few minutes we pull in posts that are getting a lot of attention
|
2. **MEASURE** We re-check each post's view count over and over, to see if it's still growing
|
3. **JUDGE** AI reads the fastest-growing ones and asks: could this be a memecoin?

- | | |
|----------|---|
| | |
| 4. GROUP | Posts about the same thing get bundled into one "narrative" |
| | |
| 5. SHOW | The best narratives appear in the app, newest and fastest first |
| | |
| 6. ACT | The trader creates the coin, or buys the one that exists |

A real example

A Norwegian ferry captain retires. On his last day he blows the ship's horn, and every other boat in the harbour blows theirs back. Someone films it and posts it.

- **9:00am** — the post goes up. Nobody in crypto is looking.
- **9:04am** — our collector picks it up. 41,000 views.
- **9:09am** — we check again: 118,000 views. It's growing fast. That growth is the signal.
- **9:12am** — someone on TikTok posts their own clip of the same moment. Now it's a *narrative*, not one lucky post.
- **9:13am** — it appears in Insidor. No coin exists. Anyone looking can be first.
- **9:31am** — a user creates \$FERRY straight from our app.
- **9:47am** — crypto Twitter finally notices. Everyone else's tools light up. Our user is already 34 minutes ahead.

3 · The machinery — what each part does

Insidor is made of six parts. Here's each one, what it does, and why it exists.

1 The Collector

Runs every couple of minutes, around the clock. It asks a data company (a "reseller") for social media posts that are getting a lot of attention right now. We can't check every post on the internet — there are hundreds of millions a day — so it filters hard: only posts above a certain view count, only recent ones.

Why it matters: if this catches the wrong things, everything downstream is junk. Everything else depends on this being good.

2 The Measurer

Once we've found a post, we keep re-checking it — every couple of minutes for the first few hours — and write down its view count each time.

Why it matters: a single view count tells you almost nothing. 500,000 views could mean it exploded in ten minutes (get in now) or crawled there over two days (it's over). Only by measuring the *same post twice* can we tell whether it's still climbing. This is the difference between us and a simple keyword alert, and it's why we spend money re-reading posts we've already seen.

Important: this history can't be bought later. Nobody sells "what did this post look like 20 minutes ago." If we're not recording it as it happens, it's gone forever.

3 The Judge (AI)

The fastest-growing posts get read by AI, which answers one question: *could someone realistically make a memecoin out of this?* It also suggests a name and a ticker.

Why it matters: most viral posts are useless to us. Politics, sports arguments, news — huge numbers, nobody makes a coin from them. What works is one absurd thing with a name: a frog called Gerald, a plastic plant, a ferry horn. The AI's job is to tell those apart. We only ask the AI about posts that are already growing fast, which keeps the cost tiny.

4 The Grouper

When something goes viral, dozens of people post about it. The Grouper spots that ten different posts are all about the same ferry captain, and bundles them into **one narrative**.

Why it matters: a feed of individual posts would show the same story ten times. Grouping also gives us a much stronger signal — one person posting is luck; twenty people independently posting about the same thing is a meme actually forming, and that's what turns into a coin.

5 The Library (database)

One central store where everything lives: every narrative, every post, every view-count reading, every coin we've matched. The collector, measurer, judge and grouper all write into it. The app reads out of it.

Why it matters: it's the single source of truth. Nothing else keeps its own copy of anything, so nothing can disagree.

6 The App

What the user actually sees and touches. It reads narratives out of the Library and shows them, and it connects to outside services to do the money parts — checking a coin is safe, showing a price chart, doing a buy, creating a new coin.

Where these live

The Collector, Measurer, Judge and Grouper are automatic jobs that run on a schedule — nobody has to press anything, and they keep running whether we're awake or not, whether our laptops are on or off. The Library is a hosted database. The App is a website.

Things we don't build ourselves

Some jobs are already solved by other companies, and building our own version would take months for no advantage. We plug into theirs:

Job	Who does it	Why not build it
Reading social media posts	A data reseller	Direct access costs tens of thousands a month
Judging "is this memeable"	Claude (AI)	Understanding humour and language is exactly what AI is for
Buying and selling coins	Axiom or Photon	Trade execution is hard, risky, and already perfect elsewhere. We use theirs inside our page.
Creating new coins	pump.fun	It's the standard everyone already uses
Connecting a wallet	Privy	Handling people's money securely is a specialist job
Prices, charts, safety checks	DexScreener, Helius, RugCheck	Standard data, freely available

What we *do* build ourselves is the part nobody else has: finding the viral post before the coin exists, and grouping it into a narrative. That's the company. Everything else is plumbing we rent.

4 · The features, one by one

Feature 1 — The Narratives feed

What it is: the home screen. A list of viral stories happening right now, best at the top.

What the user sees on each row:

- **A picture** from the post
- **A title** — e.g. "Harbor Last Horn"
- **Two lines explaining it in plain English** — "A retiring ferry captain blew his horn one last time, and every boat in the harbour blew back." You should understand the story without watching the video. This is the most important text in the whole product.
- **How many views** — big and clear, the headline number
- **A tiny line graph** showing whether views are still climbing or have flattened off
- **How old it is**
- **One button** — "Create coin" if none exists, "Buy \$TICKER" if one does, or "3 coins →" if several do

What the user does *not* see: any of our internal scoring. Behind the scenes we calculate numbers to decide what's worth showing and in what order — but we never put those on screen. The app's job is to make the judgement so the user doesn't have to. Showing our workings would just make them do our job for us.

Feature 2 — The Narrative page (the main attraction)

What it is: tap any row and you land here. This is the screen the whole product is built around. It answers "what is this, is it still hot, which coin is the real one, and what do other traders think" — all in one place, so the user never has to open another app.

What's on it:

- **The story** — the picture, the plain-English explanation, a link to the original post
- **How viral it is** — views, reposts, likes, age, and whether it's still climbing
- **Every coin made from this narrative** (see below)
- **Trading** — the chart and buy/sell built into the page, powered by Axiom or Photon
- **Research** — is the coin safe, who's holding it, who's buying, all the posts in this narrative
- **Thesis** — what other traders are saying (see below)

Why we list every coin, not just one

When something goes viral, several people often make a coin from it within minutes — one is the real one, the rest are copycats riding it. If we showed only "the coin," a user could easily buy a copycat and lose money.

So we show all of them, and label them honestly:

- **FIRST** — this was made from this narrative before any other
- **BIGGEST** — this one has the most money in it
- **UNVERIFIED** — we *think* this relates to the narrative but we're not confident, so **we don't show a buy button on it at all**

That last rule matters. If we're not sure, we don't offer the trade. A greyed-out button looks like something we're holding back; no button at all is us not making a claim we can't stand behind.

The thesis / comments layer

Traders can post their take on a narrative — "the meme isn't the captain, it's all the other boats answering." Two facts get permanently attached to every comment at the moment it's posted:

- **How big the story was when they said it** — so someone who called it at 340K views reads very differently from someone who showed up at 3 million
- **Whether they own the coin** — holding, sold, or no position

This is something no competitor can do. Twitter knows what people say but not what they own. A trading app knows what they own but not what they think. We have both — and over time, the traders with a good public record become a reason to use Insider instead of anything else.

Feature 3 — Create a coin

What it is: when a narrative has no coin yet, the user can make one in a few taps. This is our biggest differentiator — nobody else surfaces the moment *before* a coin exists.

What happens: we pre-fill the name, ticker and image from the narrative, check the ticker isn't already taken, show the cost (about 1 SOL including an optional first buy), and let the user sign it with their own wallet. The coin is created on pump.fun. The user earns a share of every future trade as the creator; Insider takes a smaller share.

Safety note: we never hold anyone's money or keys. We prepare the transaction; the user approves it in their own wallet.

Feature 4 — Trading

What it is: buying and selling, built into the narrative page so nobody has to leave.

How: we use Axiom's or Photon's trading service (whichever is cheaper) rather than building our own. Their systems already handle prices, routing and execution reliably. We earn a small fee on trades that go through us.

Feature 5 — Watchlist

Save a narrative or a coin. If a saved narrative gets its first coin, we tell the user immediately. This is the thing that brings people back — instead of them remembering to check, we open the app for them at the one moment there's something to do.

Feature 6 — Search

One box. Type a keyword, a ticker, or a coin address. Get back both the posts and the coins. Its main job is answering "does a coin already exist for this thing I just saw?" — which stops someone accidentally creating a duplicate.

Feature 7 — Trending

Rankings. Top narratives and top coins, sortable by views and age, so a user can browse rather than only seeing what we push at them.

5 · What we store

Everything lives in one database. In plain terms:

What	In plain English
Narratives	One entry per viral story — its title, plain explanation, picture, total views, age
Posts	The individual social media posts that make up each story
View readings	Every time we re-checked a post and what the view count was. This is what draws the climbing/flat line.
Coins	Every coin we've linked to a narrative, and how confident we are in that link
Comments	Trader theses, with the "how big was it when they said it" stamp frozen in
Running costs	How much we've spent today on data and AI, so we can't overspend by accident

We also calculate internal scores — how fast something is growing, how memeable the AI thinks it is. These are stored and used to decide what appears and in what order. **They are never shown to users and never sent to their browser.**

6 · The build plan

Eight stages. Each one has to work before the next makes sense.

STAGE 0 · FIRST, BEFORE ANYTHING ELSE

Prove it actually works, and tidy the foundations

The problem: right now we believe our system finds good viral posts early. We haven't proved it. Every setting we've chosen — how many views is enough, how memeable is memeable — was an educated guess.

What we do: work backwards. Take coins that already did well over the last few days, trace each one back to the viral post it came from, and check: did our system spot those posts? What did they have in common? Then compare them against posts that went viral but never became anything.

Why it matters: this either proves the whole idea is real, or tells us exactly what to fix — before we spend weeks building on top of it. It also creates the dataset behind the single most valuable thing we could eventually show a trader: *"stories like this one, at this size, historically became coins that did X."*

At the same time: the code needs organising. It grew fast and it's tangled, which makes every bug take hours instead of minutes. Mercury handles this.

STAGE 1

The narratives feed

The home screen, done properly: picture, title, the two-line plain explanation, views, the climbing line, age, one button. **Done when** a trader can scan the list and immediately understand each story without opening anything.

STAGE 2 · THE BIG ONE

The narrative page

The main attraction. Its own page with its own web address, so it can be shared into a group chat and opened by someone who's never heard of us. Includes the story, the virality numbers, every coin found for it, and the research. **Done when** a trader can decide whether to act without leaving the page.

STAGE 3

Trading built into the page

Plug in Axiom or Photon so buying happens right there. **First we need to confirm with them** that they allow this and what their fees are on top of ours. If not, we already have a working alternative.

STAGE 4

Create a coin

Wire up pump.fun properly so a user can deploy in a few taps with the name and image pre-filled. **Done when** someone can go from seeing a story to owning the first coin of it in under a minute.

STAGE 5

The thesis layer

Comments on narratives, stamped with how big the story was and whether the person holds. This is what makes people come back and what makes Insidor hard to copy. It comes after the core works, because comments on a feed nobody trusts are worthless.

STAGE 6

Search, watchlist, trending

The supporting screens. Find things, save things, browse rankings, and get told the moment a saved story gets a coin.

STAGE 7

"What happened to stories like this one"

The payoff from Stage 0. Once we have enough history, each story can show what similar ones did after they became coins. This is the thing that turns the feed from interesting into genuinely decisive — and it only exists because we've been recording since day one.

7 · What it costs to run

What	Roughly	Notes
Reading social media posts	~£4/day	By far the biggest cost. Capped so it can't overspend.
AI judging	~£1/day	Only the fastest-growing posts get read
Database, hosting, price data	Free for now	Free tiers cover us at this size
Total	~£5/day (~£150/month)	Scales up with users and coverage

Hard spending limits are built in. If we hit the daily cap, collection pauses until the next day rather than quietly running up a bill.

8 · The one thing that has to be true

BE HONEST ABOUT THIS INTERNALLY

Every feature in this document assumes one thing: that our system genuinely finds *memeable* stories *early*. Everything else — the beautiful page, the coin list, the comments — is worthless if what we surface is noise.

That's why Stage 0 comes first, and why we shouldn't build deep until it's answered. The good news: we can answer it in days, not months, and either outcome is useful. If the signal is real, we have proof to show investors. If it isn't, we find out before wasting a month.

Glossary

Term	Means
Narrative	A viral story — one or more posts about the same thing, bundled together
Ticker	A coin's short name, like \$FERRY
Mint / contract address	A coin's unique ID on the blockchain
pump.fun	The site most Solana memecoins are created on
Wallet	Where a user keeps their crypto. They approve every transaction themselves.
Liquidity	How much money is in a coin — low liquidity means you may not be able to sell
Market cap (mcap)	What the whole coin is worth right now
Rug	When a coin's creator takes the money and the coin collapses
CT	Crypto Twitter — the crowd we're trying to be ahead of
Trencher	A trader who hunts new memecoins all day. Our user.

Insidor · we find the viral story before the coin exists · feed → narrative page → create or buy · we build the detection, we rent everything else ·
prove the signal before building deep.